

Kokkos 4.1 Release Briefing

New Capabilities

July 20, 2023

4.1 Release Highlights

- ▶ Backend updates
- ▶ Build system updates
- ▶ Multiple Reducers for Nested Parallel Reduce
- ▶ Bit Manipulation
- ▶ UnorderedMap Insertion Operation Types
- ▶ Miscellaneous
- ▶ Deprecations and other breaking changes
- ▶ Kokkos on Compiler Explorer

Online Resources:

- ▶ <https://github.com/kokkos>:
 - ▶ Primary Kokkos GitHub Organization
- ▶ <https://github.com/kokkos/kokkos-tutorials/wiki/Kokkos-Lecture-Series>:
 - ▶ Slides, recording and Q&A for the Full Lectures
- ▶ <https://kokkos.github.io/kokkos-core-wiki>:
 - ▶ Wiki including API reference
- ▶ <https://kokkosteam.slack.com>:
 - ▶ Slack channel for Kokkos.
 - ▶ Please join: fastest way to get your questions answered.
 - ▶ Can whitelist domains, or invite individual people.


Would like to strengthen community bonds and discoverability

List of Applications and Libraries

- ▶ Add your app to <https://github.com/kokkos/kokkos/issues/1950>
- ▶ We are planning to add that to a Kokkos website.
- ▶ Helps people discover each other when working on similar things.

GitHub Topics

- ▶ Use *kokkos* tag on your repos.
- ▶ If you click on the topic you get a list of all projects on github with that topic.



Explore **Topics** Trending Collections Events GitHub Sponsors

Get email updates

kokkos

★ Unstar

Here are 32 public repositories matching this topic...

Language: All ▾

Sort: Best match ▾

📁 kokkos / kokkos

★ Starred 1.1k ▾

<> Code Issues Pull requests

Kokkos C++ Performance Portability Programming EcoSystem: The Programming Model - Parallel Execution and Memory Abstraction

c-plus-plus parallel-computing abstraction high-performance-computing programming-model

kokkos sni-prog-models-runtimes

Updated 1 minute ago C++

📁 ParRes / Kernels

☆ Star 347 ▾

<> Code Issues Pull requests

This is a set of simple programs that can be used to explore the features of a parallel platform.

Improve this page

Add a description, image, and links to the kokkos topic page so that developers can more easily learn about it.

[Curate this topic](#) >

Add this topic to your repo

To associate your repository with the kokkos topic, visit your repo's landing page and select "manage topics."

[Learn more](#) >

- ▶ We are considering organizing a multi-day in-person user group meeting
- ▶ Likely in Albuquerque
- ▶ postponed to early December time frame
- ▶ Tentative content
 - ▶ Updates from the Kokkos team (new features and planned work)
 - ▶ User experiences: porting to AMD and Intel GPUs
 - ▶ User experiences: performance portability studies
 - ▶ Best practices (also user-provided)
 - ▶ Students and Postdocs showcase
 - ▶ Feedback and discussion session

Backend Updates

Content:

- ▶ Backend Updates I
- ▶ Backend Updates II

CUDA

- ▶ Allow NVCC 12 to compile using C++20 flag
- ▶ Remove ability to disable CMake option `Kokkos_ENABLE_CUDA_LAMBDA` and unconditionally enable CUDA extended lambda support.
- ▶ Drop unnecessary fences around the memory allocation when using `CudaUVMSpace` in views
 - ▶ Issues when relying on View initialization to fence when a execution space instance was passed.

HIP

- ▶ Improve performance for `parallel_reduce`. Use different parameters for `LightWeight` kernels.

SYCL

- ▶ Improve and simplify `parallel_scan` implementation

OpenMPTarget

- ▶ Improve hierarchical parallelism for Intel architectures
- ▶ Enable Cray compiler for the `OpenMPTarget` backend.

HPX

- ▶ Update HPX backend to use HPX's sender/receiver functionality
- ▶ Increase minimum required HPX version to 1.8.0
- ▶ Implement `HPX::in_parallel`

- ▶ Export CMake `Kokkos_CUDA`, `HIP_ARCHITECTURES` variables
- ▶ Allow linking against build tree
- ▶ Kokkos can be used as an external dependency in Trilinos
- ▶ Drop `Kokkos_ENABLE_LAUNCH_COMPILER` option which had no effect
- ▶ Export variables for relevant Kokkos options with CMake

Multiple Reducers for Nested Parallel Reduce

Content: Team-level parallel reduce with multiple reducers

- ▶ Extended reducer capabilities in nested `parallel_reduce`
- ▶ Allow multiple reductions in a single team `parallel_reduce`
- ▶ Supported for `TeamThreadRange`, `ThreadVectorRange` and `TeamVectorRange` policies
 - ▶ Not available for `TeamMDRangePolicies` for now

```
template <typename TeamPolicy, typename FunctorType,  
          typename... ReducerArgument>  
Kokkos::parallel_reduce(const TeamPolicy& policy,  
                       const FunctorType& functor,  
                       const ReducerArgument&... reducer);  
  
template <typename TeamPolicy, typename FunctorType,  
          typename... ReducerArgumentNonConst>  
Kokkos::parallel_reduce(const TeamPolicy& policy,  
                       const FunctorType& functor,  
                       ReducerArgumentNonConst&... reducer);
```

- ▶ The number of reducers and the number of functor's reducer value arguments must match.

```
Kokkos::parallel_for(
  policy, KOKKOS_LAMBDA(team_member_type const& team) {
    /* ... */

    Kokkos::parallel_reduce(
      teamPolicy,
      [=](int& i, int& arg0, int& arg1, int& arg2, int& arg3) {
        /* ... */
      },
      result0, Kokkos::Prod<int>(result1),
      Kokkos::Max<int>(result2), result3);
  }
);
```

Bit Manipulation

Content:

- ▶ Kokkos:: equivalents for C++20/C++23 components to access, manipulate and process both individual bits and bit sequences
 - ▶ bit_cast
 - ▶ byteswap
 - ▶ Integral powers of 2
 - ▶ has_single_bit, bit_ceil, bit_floor, bit_width
 - ▶ Rotating
 - ▶ rotl, rotr
 - ▶ Counting
 - ▶ countl_zero, countl_one, countr_zero, countr_one, popcount

~~constexpr~~ To `bit_cast<To>(From const& from)` noexcept

- ▶ Reinterpret the object representation of one type as that of another
 - ▶ `sizeof(From) == sizeof(To)`
 - ▶ From must be trivially copyable
 - ▶ To must be trivially copyable
- ▶ Not `constexpr` (differs from C++23 `std::bit_cast`)

```
double d1 = 19880124.0;
auto u64 = Kokkos::bit_cast<uint64_t>(d1);
auto d2 = Kokkos::bit_cast<double>(u64);

assert(d1 == d2);
```

`constexpr T byteswap(T value) noexcept`

- ▶ Reverses the bytes in the given integer value

- ▶ T is an integral type
 - ▶ `bool`, `char`, `char8_t`, `char16_t`, `char32_t`, `wchar_t`, `short`, `int`, `long`, `long long`, `clang __128` (but not `gcc __128`)
 - ▶ signed and unsigned integer types

```
int32_t i1 = 0xdeadbeef;  
auto    i2 = Kokkos::byteswap(i1);  
  
assert(i2 == 0xefbeadde);
```


- ▶ `constexpr bool has_single_bit(T x) noexcept`
 - ▶ Checks if a number is an integral power of 2
- ▶ `constexpr T bit_ceil(T x) noexcept`
 - ▶ Finds the smallest integral power of two not less than `x`
- ▶ `constexpr T bit_floor(T x) noexcept`
 - ▶ Finds the largest integral power of 2 not greater than `x`
- ▶ `constexpr int bit_width(T x) noexcept`
 - ▶ Finds the smallest number of bits needed to represent `x`

- ▶ `T` is an unsigned integer type
 - ▶ `unsigned char`, `unsigned short`, `unsigned int`,
`unsigned long`, `unsigned long long`

```
uint64_t x = 5; // 0b101
assert(Kokkos::has_single_bit(x) == false);
assert(Kokkos::bit_ceil(x) == 8);
assert(Kokkos::bit_floor(x) == 4);
assert(Kokkos::bit_width(x) == 3);
```

- ▶ `constexpr T rotl(T x, int x) noexcept`
 - ▶ Computes the result of a bitwise left-rotation
- ▶ `constexpr T rotr(T x, int x) noexcept`
 - ▶ Computes the result of a bitwise right-rotation

- ▶ T is an unsigned integer type

```
uint16_t i16 = 0b1001110000111001;  
assert(Kokkos::rotl(i16, 2) == 0b0111000011100110);  
assert(Kokkos::rotr(i16, 3) == 0b0011001110000111);
```

- ▶ `constexpr int countl_zero(T x) noexcept`
 - ▶ Count the number of consecutive 0 bits, starting from the most significant bit
- ▶ `constexpr int countl_one(T x) noexcept`
 - ▶ Count the number of consecutive 1 bits, starting from the most significant bit
- ▶ `constexpr int countr_zero(T x) noexcept`
 - ▶ Count the number of consecutive 0 bits, starting from the least significant bit
- ▶ `constexpr int countr_one(T x) noexcept`
 - ▶ Count the number of consecutive 1 bits, starting from the least significant bit
- ▶ `constexpr int popcount(T x) noexcept`
 - ▶ Count the number of 1 bits in an unsigned integer

- ▶ T is an unsigned integer type

```
uint16_t bits = 0b1111101000110100;  
  
assert(Kokkos::countl_zero(bits) == 0);  
assert(Kokkos::countl_one(bits)  == 5);  
assert(Kokkos::countr_zero(bits) == 2);  
assert(Kokkos::countr_one(bits)  == 0);  
assert(Kokkos::popcount(bits)    == 9);
```

- ▶ In namespace `Kokkos::Experimental::`
- ▶ Not `constexpr`
- ▶ Directly call the compiler builtin version, if beneficial

- ▶ `bit_cast_builtin`
- ▶ `byteswap_builtin`
- ▶ Integral powers of 2
 - ▶ `has_single_bit_builtin`, `bit_ceil_builtin`,
`bit_floor_builtin`, `bit_width_builtin`
- ▶ Rotating
 - ▶ `rotr_builtin`, `rotr_builtin`
- ▶ Counting
 - ▶ `countl_zero_builtin`, `countl_one_builtin`,
`countr_zero_builtin`, `countr_one_builtin`,
`popcount_builtin`

UnorderedMap Insertion Operation Types

Content: Extended UnorderedMap insertion behavior

- ▶ Default behavior is to insert a key, value pair exactly once
- ▶ Maintain default behavior via operation type `NoOp`
- ▶ Allow existing key, value pairs to be accumulated into via operation type `AtomicAdd`

```
template <class ValueTypeView, class ValuesIdxType>
struct UnorderedMapInsertOpTypes {
    using value_type = typename ValueTypeView::non_const_value_type;
    struct NoOp {
        void op(ValueTypeView, ValuesIdxType, const value_type);
    };
    struct AtomicAdd {
        void op(ValueTypeView values, ValuesIdxType values_idx,
               const value_type v);
    };
};
```

```
template <typename InsertOpType = default_op_type>
insert_result insert(key_type const &key,
                    impl_value_type const &value,
                    InsertOpType arg_insert_op = InsertOpType());
```

- ▶ For other use-cases, more operation types can be added to `UnorderedMapInsertOpTypes`

```
using map_op_type
    = Kokkos::UnorderedMapInsertOpTypes<value_view_type, size_type>;
using atomic_add_type = typename map_op_type::AtomicAdd;
atomic_add_type atomic_add;
parallel_for(N, KOKKOS_LAMBDA (uint32_t i) {
    map.insert(i, values(i), atomic_add);
});
```


▶ Add `Kokkos::Profiling::ScopedRegion`

```
double myfunction()
{
    Kokkos::Profiling::ScopedRegion region("foo");
    if (...)
        return bar;
    else
        return eval();
}
```

- ▶ Add support for `View::rank[_dynamic]()`
- ▶ Detect incompatible relocatable device code mode to prevent ODR violations
- ▶ Add (experimental) support for 32-bit Darwin and PPC
- ▶ Add missing half and bhalf specialization of the infinity numeric trait

- ▶ Add `is_dual_view` trait and align template parameters with regular view
- ▶ Allow templated functors in `parallel_for`, `parallel_reduce`, and `parallel_scan`
- ▶ Define `KOKKOS_COMPILER_INTEL_LLVM` and only define at most one `KOKKOS_COMPILER*` macro

Deprecations and Behavior Changes

Content:

- ▶ Legacy Atomics Removal
- ▶ CUDA_LAMBDA always on
- ▶ Fencing and UVM
- ▶ Other

Legacy Atomics Fallback removed

DESUL atomics have been the default for a few releases, now removed option to use legacy atomics

Related are a few macro removals:

- ▶ `KOKKOS_ENABLE_CUDA_ASM*` - macros only used in legacy atomic implementation
- ▶ `KOKKOS_ENABLE_[CUDA/OPENMP/GNU/INTEL]_ATOMICS` - used to define legacy atomics

- ▶ Removed option to enable/disable CUDA lambdas - they are always on
- ▶ LAMBDA support is stable in enough in all supported compilers

Configure behavior depends on deprecation setting:

- ▶ `Kokkos_ENABLE_CUDA_LAMBDA` not set: **recommended!**
- ▶ `Kokkos_ENABLE_CUDA_LAMBDA=ON`: ok now, no warning - future release unused variable warning
- ▶ `Kokkos_ENABLE_CUDA_LAMBDA=OFF`
 - ▶ `Kokkos_ENABLE_DEPRECATED_CODE_4` not set: warning about setting lambda options, and that the setting is ignored
 - ▶ `Kokkos_ENABLE_DEPRECATED_CODE_4=ON`: warning about setting lambda options, and that the setting is ignored
 - ▶ `Kokkos_ENABLE_DEPRECATED_CODE_4=OFF`: configure error

- ▶ Continue to work on removing fencing which isn't semantically required by Kokkos
- ▶ Now removed fence for CudaUVMSpace View creation when providing an execution space instance

```
// No change in behavior - implicit fence after init  
View<int*, CudaUVMSpace> a("A", N);  
a[0] = 3;
```

```
// Relaxed fencing in 4.1 to implement desired async behavior  
View<int*, CudaUVMSpace> a(view_alloc(Cuda(), "A"), N);  
// required fence in 4.1 to prevent race condition below  
Cuda().fence();  
a[0] = 3;
```

General Rule:

Kokkos operations taking an execution space instance are asynchronous!

Trilinos/TriBITS related:

- ▶ Removed TriBITS Kokkos subpackages - so if you get Kokkos from Trilinos, not more kokkos-core/algorithms/containers separation
- ▶ Removed associated (unused) Kokkos[Algorithms,Containers]_config.h files - we did not find any project including them

Sorting:

- ▶ Removed default constructors of BinSort, BinOp1D, and BinOp3D - the default constructed state was invalid and can lead to hard to understand failures.

How to Get Your Fixes and Features into Kokkos

- ▶ Fork the Kokkos repo (<https://github.com/kokkos/kokkos>)
- ▶ Make topic branch from *develop* for your code
- ▶ Add tests for your code
- ▶ Create a Pull Request (PR) on the main project *develop*
- ▶ Update the documentation (<https://github.com/kokkos/kokkos-core-wiki>) if your code changes the API
- ▶ Get in touch if you have any questions (<https://kokkosteam.slack.com>)