

# Kokkos 4.6 Release Briefing

New Capabilities

04/10/2025

## 4.6 Release Highlights

- ▶ Organizational
- ▶ Feature Highlights
- ▶ General Enhancements
- ▶ Backend updates
- ▶ Build system updates
- ▶ Deprecations and other breaking changes
- ▶ Bug Fixes

## Online Resources:

- ▶ <https://github.com/kokkos>:
  - ▶ Primary Kokkos GitHub Organization
- ▶ <https://kokkos.org/kokkos-core-wiki/tutorials-and-examples.html>:
  - ▶ Tutorials, video lectures, and examples
- ▶ <https://kokkos.org/kokkos-core-wiki>:
  - ▶ Wiki including API reference
- ▶ <https://kokkosteam.slack.com>:
  - ▶ Slack workspace for Kokkos.
  - ▶ Please join: fastest way to get your questions answered.
  - ▶ Can whitelist domains, or invite individual people.

## Would like to strengthen community bonds and discoverability

### *List of Applications and Libraries*

- ▶ Add your app to <https://github.com/kokkos/kokkos/issues/1950>
- ▶ We are planning to add that to the Kokkos website.
- ▶ Helps people discover each other when working on similar things.

### *GitHub Topics*

- ▶ Use *kokkos* tag on your repos.
- ▶ If you click on the topic you get a list of all projects on github with that topic.

# Organizational

## Content:

- ▶ HPSF and Kokkos Meeting 2025
- ▶ Targeting C++20 for Kokkos 5.0
- ▶ Makefile deprecation

## Kokkos User Group Meeting 2025 @ HPSF Conference

- ▶ *When:* May 5th-8th 2025
- ▶ *Where:* Chicago
- ▶ *What:* 2-days HPSF plenary + 2-days Project meetings
- ▶ *KUG-Content:* Focused on user experiences
  - ▶ How do you leverage Kokkos?
  - ▶ What are pain points?
  - ▶ Kokkos-based libraries of interest to the community

*Registration open now!*

## What to expect from KUG

- ▶ Eight 90-minute sessions featuring a dynamic blend of Kokkos developers and community users
- ▶ *Day 1 Highlights:*
  - ▶ Essential Updates
  - ▶ Kokkos in Applications
  - ▶ Adopting Kokkos
  - ▶ Lightning Talks
- ▶ *Day 2 Highlights:*
  - ▶ Kokkos Ecosystem
  - ▶ Tuning and Performance
  - ▶ Algorithms
  - ▶ Panel Discussion

## Other reasons to go

- ▶ General Poster Session
- ▶ Updates on the HPSF project
- ▶ Introduction to various working groups
- ▶ Various Panel Discussions
- ▶ Chance to meet all other members of HPSF
- ▶ ...



- ▶ HPSF will be present at [ISC BOF 2025](#)
- ▶ [Kokkos Tea-Time](#) on 2nd or 3rd Wed of the month
  - ▶ April 16th @ 11am EST "Solomon: unified schemes for directive-based GPU offloading"

**Kokkos 5 is coming Summer 2025**

**We will require C++20!**

*Start preparing now:*

- ▶ Check availability of compilers on your systems
- ▶ Test with C++20 enabled: start with a CPU build
- ▶ Minimum Compiler requirements will change (more details later)

*Nothing wrong for your project to require C++20 now if you feel ready!*

**Makefile is officially deprecated and will be removed in the next major release**

*Start preparing now:*

- ▶ Check if you can transition to CMake
- ▶ Comment on pinned issue [7610](#)

## We reached “passed” on the OSSF Best Practices Program

[www.bestpractices.dev](http://www.bestpractices.dev)

*This means Kokkos is continuously tracking and openly reporting the conformity with open source software practices.*

# Feature highlights

- ▶ describes asynchronous workloads organised as a direct acyclic graph (DAG)
- ▶ executed using `submit()`, possibly many times, observing dependencies

```
auto graph = Kokkos::create_graph([&](auto root) {  
    auto node_A = root.then_parallel_for("A", ...policy..., ...functor...);  
    auto node_B = node_A.then_parallel_for("B", ...policy..., ...functor...);  
    auto node_C = node_A.then_parallel_for("C", ...policy..., ...functor...);  
  
    auto node_D = Kokkos::when_all(node_B, node_C).  
        then_parallel_for("D", ...policy..., ...functor...);  
});  
  
graph.instantiate();  
  
graph.submit();
```

- ▶ `then` node: executes a callable on device
- ▶ Single call of the functor per `submit()`
- ▶ Executed in the `ExecutionSpace` the graph is submitted to

```
auto graph = Kokkos::create_graph([&](auto root) {  
    auto node_A = root.then_parallel_for("A", ...policy..., ...functor...);  
    auto node_B = node_A.then("B", ...functor...);  
});
```

- ▶ Functor passed to `then` must be callable without arguments and marked with `KOKKOS_FUNCTION`

- ▶ Interoperability: create a `Kokkos::Graph` from a native Cuda/HIP/Sycl graph

```
cudaGraph_t native_graph = nullptr;  
cudaGraphCreate(&native_graph, 0);
```

```
auto graph_from_native =  
    Kokkos::Experimental::create_graph_from_native(exec, native_graph);
```

- ▶ Experimental, does not yet allow adding nodes created using the native API to a `Kokkos::Graph`



- ▶ Launch kernels on multiple devices from a single host process
- ▶ Available for ROCm 5.6 and later
- ▶ Requires direct use of HIP runtime API for creating and destroying streams
- ▶ Experimental, still looking for feedback from new users
- ▶ New documentation (for all backends)

<https://kokkos.org/kokkos-core-wiki/API/core/MultiGPUSupport.html>

```
// Create streams on different devices
hipStream_t streams[2];
hipSetDevice(0); hipStreamCreate(&streams[0]);
hipSetDevice(1); hipStreamCreate(&streams[1]);
{
    // Creating execution spaces
    Kokkos::HIP exec0(streams[0]), exec1(streams[1]);

    // Allocating views
    Kokkos::View<int*> v0(Kokkos::view_alloc("v0", exec0), N);
    Kokkos::View<int*> v1(Kokkos::view_alloc("v1", exec1), M);

    // Launch kernels (run concurrently)
    Kokkos::parallel_for(Kokkos::RangePolicy(exec0, 0, N), functor0);
    Kokkos::parallel_for(Kokkos::RangePolicy(exec1, 0, M), functor1);
}
// Destroy streams (after execution spaces are deleted)
hipStreamDestroy(streams[0]); hipStreamDestroy(streams[1]);
```

# General Enhancements

- ▶ `kokkos_check`: Check at configure time that Kokkos was built with the requested backends and target architectures.
- ▶ Fix a warning when a user calls the cmake function `kokkos_check` from a `<PackageName>Config.cmake` "Find Module" file

```
CMake Warning (dev) at /usr/share/cmake-3.22/Modules/FindPackageHandleStandardArgs.cmake:438 (message):  
  The package name passed to 'find_package_handle_standard_args'  
  (Kokkos_DEVICES) does not match the name of the calling package (SomePackage).  
  This can lead to problems in calling code that expects 'find_package'  
  result variables (e.g., '_FOUND') to follow a certain pattern.  
Call Stack (most recent call first):  
  ... /kokkos/lib/cmake/Kokkos/KokkosConfigCommon.cmake:110 (find_package_handle_standard_args)  
  ...  
This warning is for project developers. Use -Wno-dev to suppress it.
```

With the Cuda and HIP backends, `Kokkos::Experimental::inclusive_scan` now calls the vendor versions in Thrust

- ▶ The vendor versions are up to 3x faster than the `Kokkos::parallel_scan`-based default implementation
- ▶ Thrust requires `Kokkos_ENABLE_ROCTHRUST` to be ON (which is the default)
- ▶ Approximately 1.5-3x speed up (V100, MI300A)

Reduced the overhead of Kokkos tools related checks

- ▶ Store the information whether Kokkos tools are enabled after each modification to the tools' callbacks
- ▶ Previously, this value was recomputed for every event (`parallel_for`, `fence`, etc.)
- ▶ Most noticeable in small serial kernels (around 100 elements)
- ▶ Reduction of launch time of approximately 10ns (About the time to increment 100 elements in a kernel) on CPU

Added reductions and remaining compound assignments to Kokkos SIMD


- ▶ `basic_simd& operator/=(basic_simd&, U&&)`
- ▶ `basic_simd& operator>>=(basic_simd&, U&&)`
- ▶ `basic_simd& operator<<=(basic_simd&, U&&)`
  
- ▶ `T reduce_min(const basic_simd& x)`
- ▶ `T reduce_max(const basic_simd& x)`
- ▶ `T reduce(const basic_simd& x, const mask_type& mask, T identity_element, BinaryOperation binary_op)`
  - ▶ Supported binary operations are: `std::plus`, `std::multiplies`, `std::bit_and`, `std::bit_or` and `std::bit_xor`
  - ▶ `std::plus` is used if binary op is not specified

- ▶ In Kokkos 4.5, we fixed a performance bug in `Kokkos::sort`
- ▶ Root cause is in an implementation detail in `RandomAccessIterator`
  
- ▶ In Kokkos 4.6 the root cause is fixed
- ▶ Fixes all algorithms that rely on `RandomAccessIterator` (e.g. `sort`, `search`, etc.)



**Improvement dependent on algorithm and hardware!**

Host only, Intel Skylake



`print_configuration` outputs if system allocated memory is accessible on GPU

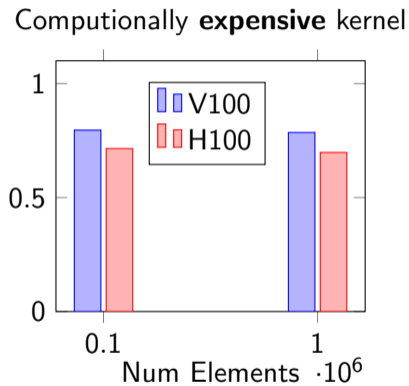
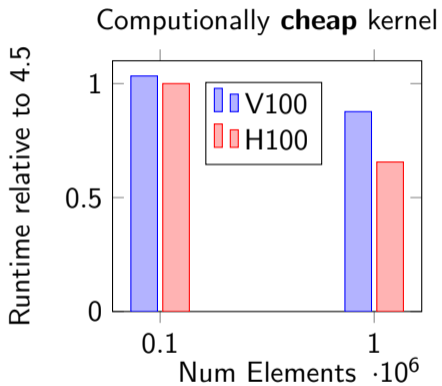
**No guarantees about the print format!**

Example output for MI300A with HIP backend

```
XNACK environment variable set: yes
Kernel reports HMM module via 'CONFIG_HMM_MIRROR=y' in '/boot/config': yes
Architecture capable of accessing system allocated memory: 1,
System allows accessing system allocated memory on GPU: 1,
```

# Backend Updates

- ▶ SYCL: Improved sorting performance for non-contiguous views
- ▶ Serial: Reduce fences overhead when using `Kokkos_ENABLE_ATOMICS_BYPASS`
- ▶ CUDA: Improved performance for `Kokkos::parallel_reduce` on H100 and newer by removing limitations on the runtime thread configuration



- ▶ Change block size deduction to prefer smaller blocks/teams if possible
- ▶ Allocate memory with stream ordered semantics (*i.e.* use `hipMallocAsync`)
- ▶ Fix a segfault when a virtual function called inside a kernel requires too many registers

# Build Systems Updates

- ▶ Add support for Zen 4 AMD microarchitecture (`Kokkos_ARCH_ZEN4`)
- ▶ Enable NVIDIA Grace architecture with NVHPC (`Kokkos_ARCH_ARMV9_GRACE`)
- ▶ Support static library builds via `CMAKE_CUDA_RUNTIME_LIBRARY=static` when using CUDA as CMake language



- ▶ Spack *develop* branch now supports MI300A with a new variant **apu** ([spack/spack#48609](https://github.com/spack/spack/pull/48609))
- ▶ To compile Kokkos for MI300A, forcing the APU mode, use the following command: `spack install kokkos +rocm amdgpu_target=gfx942 +apu`

# Deprecations and other breaking changes

- ▶ Intel has deprecated Intel Classic in 2022, and removed it from oneAPI 2024
- ▶ In order to focus on newer compilers, and reduce maintenance burden, we have **removed** support for Intel Classic (oneAPI Intel/icpx still supported of course!)

## Deprecate direct access to `d_view` and `h_view`

- ▶ Modifying the allocations in `d_view` and `h_view` directly is dangerous, especially if `modify` and `sync` are skipped
- ▶ Use `view_host()` and `view_device()` instead
- ▶ These two functions return by value with deprecated code enabled and by const reference otherwise. This might have performance implications if used extensively, e.g., in loop bounds.

- ▶ `native_simd`, `native_simd_mask` **deprecated** to align with the C++26 standard
- ▶ **Removed** Obtaining a reference from SIMD operator `[]` to align with the C++26 Standard
- ▶ **Changed** the return type of SIMD operator `==` and operator `!=` to return SIMD masks instead of `bool`
  - ▶ If you want old behavior, use `all_of(a == b)`

- ▶ Already discussed deprecating the Makefile
- ▶ StaticCrsGraph is **moved** to Kokkos Kernels and **deprecated** in Core
  - ▶ See <https://github.com/kokkos/kokkos-kernels/pull/2419>
  - ▶ Symbol is in Kernels under `KokkosSparse::StaticCrsGraph`

# Bug Fixes

- ▶ Fix execution of ranges with more than 2 billion elements
- ▶ Graph:
  - ▶ Fix graph node lifetime issues
  - ▶ Fix lock-based atomics failure when launching CUDA and HIP graphs
- ▶ CUDA backend: Fix incorrect iteration in MDRangePolicy of rank  $> 4$  for high iteration counts
- ▶ SIMD:
  - ▶ fix a bug in scalar min/max
  - ▶ fix a bug in non-masked reductions
- ▶ View: fix MSVC compilation



- ▶ Fix `c1lean` target when embedding Kokkos in another project
- ▶ Stop generation if ARMv9 Grace arch is not explicitly supported by the compiler when `KOKKOS_ARCH_ARMV9_GRACE` is specified
  - ▶ Can still try and configure with `ARCH_NATIVE`
- ▶ Fix Zen3 flag for NVHPC
- ▶ Use right arch for MI300A in makefiles
- ▶ (CUDA) ignore gcc assembler options in `nvcc_wrapper`

- ▶ Fix performance bug affecting `atomic_fetch_{add,sub,min,max,and,or,xor}` on integral types `long` and `unsigned long` with HIP
- ▶ Fix performance of `RangePolicy` where an error message is generated even if precondition not violated

## How to Get Your Fixes and Features into Kokkos

- ▶ Fork the Kokkos repo (<https://github.com/kokkos/kokkos>)
- ▶ Make topic branch from *develop* for your code
- ▶ Add tests for your code
- ▶ Create a pull request (PR) on the main project *develop*
- ▶ Update the documentation (<https://github.com/kokkos/kokkos-core-wiki>) if your code changes the API
- ▶ Get in touch if you have any question (<https://kokkosteam.slack.com>)