

Kokkos 5.0 Release Briefing

11/25/2025

5.0 Release Highlights

- ▶ Organizational
- ▶ Feature Highlights
- ▶ General Enhancements
- ▶ Backend updates
- ▶ Build system updates
- ▶ Deprecations and other breaking changes
- ▶ Bug Fixes

Online Resources:

- ▶ <https://github.com/kokkos>:
 - ▶ Primary Kokkos GitHub Organization
- ▶ <https://kokkos.org/kokkos-core-wiki/tutorials-and-examples.html>:
 - ▶ Tutorials, video lectures, and examples
- ▶ <https://kokkos.org/kokkos-core-wiki>:
 - ▶ Wiki including API reference
- ▶ <https://kokkosteam.slack.com>:
 - ▶ Slack workspace for Kokkos.
 - ▶ Please join: fastest way to get your questions answered.
 - ▶ Can whitelist domains, or invite individual people.

Would like to strengthen community bonds and discoverability

List of Applications and Libraries

- ▶ Add your app to <https://github.com/kokkos/kokkos/issues/1950>
- ▶ We are planning to add that to the Kokkos website.
- ▶ Helps people discover each other when working on similar things.

GitHub Topics

- ▶ Use *kokkos* tag on your repos.
- ▶ If you click on the topic you get a list of all projects on github with that topic.

Organizational

Content:

- ▶ Kokkos Training & Hackathon @ CEA
- ▶ Kokkos User Group @ HPSF Community Summit '26 (Europe)
- ▶ Kokkos User Group @ HPSF Conference '26 (US)
- ▶ Internship program
- ▶ Mailing Lists

Kokkos Training & Hackathon @ CEA, France

- ▶ When: January 12th-16th 2026 (noon to noon)
- ▶ Where: CEA, Saclay ("near" Paris), France
- ▶ What:
 - ▶ 1 day of basic training: masterclass + exercises
 - ▶ 3 days Hackathon to get your hands dirty & discover Kokkos for real:
 - ▶ get a code half-ported to GPU,
 - ▶ in small teams, with support from the devs., finish porting & optimize for GPU,
 - ▶ compete with the other teams for the best performance!



Registration is free, but mandatory, seats are limited!

<https://indico.math.cnrs.fr/e/kokkos-hackathon-26>

Kokkos User Group @ HPSF Community Summit, Europe

- ▶ *When:* February 25th-27th 2026
- ▶ *Where:* TU Braunschweig (“near” Berlin), Germany
- ▶ *What:*
 - ▶ Few HPSF plenary talks
 - ▶ Project community gathering:
Kokkos, Trilinos and WarpX user & developers groups
 - ▶ Some Spack, Module & other HPSF projects



Call for Abstracts open!

<https://events.academiccloud.de/e/hpsfcs26>

Kokkos User Group @ HPSF Conference, USA

- ▶ *When:* March 16th-20th 2026
- ▶ *Where:* Chicago, IL, USA
- ▶ *What:* 2½-days HPSF plenary + 2½-days Project meetings
- ▶ *KUG-Content:* Focused on user experiences
 - ▶ How do you leverage Kokkos?
 - ▶ What are pain points?
 - ▶ Kokkos-based libraries of interest to the community



Call for Papers & Registration open!

<https://events.linuxfoundation.org/hpsf-conference/>

SRP Internship Program

Kokkos will take part in HPSF internship program in 2026

With Sustainable Research Pathways

More info soon

Mailing Lists

Sign up for the Kokkos mailing list to stay up-to-date
with the latest Kokkos news.

<https://kokkos.org/community/mailing-lists/>



Feature highlights

Minimum compiler requirements after requiring C++20

	Kokkos 4.x	Kokkos 5.0
GCC	8.2.0	10.4.0
Clang (CPU)	8.0.0	14.0.0
Clang (CUDA)	10.0.0	15.0.0
InteLLVM (CPU)	2021.1.1	2022.0.0
InteLLVM (SYCL)	2023.0.0	2024.2.1
NVCC	11.0	12.2.0
NVC++	22.3	22.3
ROCM	5.2.0	6.2.0
MSVC	19.29	19.30

Note: Clang (CUDA) allows for CUDA 11.8 as underlying runtime.

Note: MSVC is only actually tested with the latest version.

<https://kokkos.org/kokkos-core-wiki/get-started/requirements.html>

Deprecated Code 4 is turned off by default and started to be removed!

- ▶ In Kokkos 5.0 Kokkos_ENABLE_DEPRECATED_CODE_4=OFF is the default!
- ▶ Some deprecated feature will immediately be removed.
- ▶ Features where we expect users still needing time, will stick around for a couple more minor releases.

Removed features

- ▶ Tasking
- ▶ Makefile support

- ▶ This release contains an extensive refactoring of Kokkos::View
- ▶ View was refactored to use `mdspan`, a C++23 addition to the standard library
 - ▶ `mdspan` is backported to C++17/C++20, and our implementation can be found at github.com/kokkos/mdspan/
- ▶ The goal of this refactor was to provide better library interoperability, more API flexibility, and reduced maintenance burden
- ▶ In principle, this update should be fully transparent; i.e. your existing code should work as it did before and we've done extensive testing to ensure this
- ▶ What does this mean for applications?
 - ▶ We use the same customization points as `mdspan` now, including accessors and layout mappings
 - ▶ In the future, we may provide a mechanism for users to customize these. Would that be useful for people?

Haven't we seen this before??

Yes we did: first attempt to enable in 4.7.0, then disabled in 4.7.1

- ▶ Had to disable in Trilinos (and thus Spack) anyway due to needed Trilinos refactor
 - ▶ They used lots of Impl namespace features from Kokkos
- ▶ Got reports in 4.7.0 of performance regressions
 - ▶ Addressed with workaround for implicit integer conversion when indexing into View
 - ▶ Where is a bit more pain to come to clean up 32bit vs 64bit indexing
- ▶ More issues when hoped for with older compilers
 - ▶ Most are not allowed anymore due to C++20 requirements
 - ▶ We are getting initial reports about issues with GCC 10.4 as well as some older CUDA (before 12.6) in corner cases

Recommendation

Consider updating to CUDA 12.8 and GCC 12 if possible – Clang variants are all fine.

Kokkos 5 added support for C++20 modules

- ▶ Tested with clang-19 and expected to work with host backends

Available modules

- ▶ `kokkos.core`
- ▶ `kokkos.bitset`
- ▶ `kokkos.dual_view`
- ▶ `kokkos.dyn_rank_view`
- ▶ `kokkos.dynamic_view`
- ▶ `kokkos.error_reporter`
- ▶ `kokkos.functional`
- ▶ `kokkos.offset_view`
- ▶ `kokkos.scatter_view`
- ▶ `kokkos.unordered_map`
- ▶ `kokkos.random`
- ▶ `kokkos.sort`
- ▶ `kokkos.std_algorithms`
- ▶ `kokkos.simd`

See https://github.com/kokkos/kokkos/tree/develop/example/build_cmake_installed_modules

CMake

```
# Enable default source code file scanning
cmake_minimum_required(VERSION 3.28.2)
project(Example CXX)
find_package(Kokkos REQUIRED)
add_executable(example cmake_example.cpp)
# FIXME_MODULES Use Kokkos::kokkos again
target_link_libraries(
    example
    Kokkos::kokkoscore)
```

C++

```
import kokkos.core;
#include <Kokkos_Macros.hpp>

int main(int argc, char* argv[]) {
    Kokkos::initialize(argc, argv);
    [...]
    Kokkos::finalize();
}
```

General Enhancements

- ▶ Improved first-touch handling for View initialization in Host parallel backends
 - ▶ Initialization is now done with `parallel_for` for large allocations
 - ▶ `std::memset` is used for small allocations to avoid parallel overhead
 - ▶ Threshold is chosen based on benchmarking results
- ▶ Extended UnorderedMap to support View as `value_type`
 - ▶ The `SequentialHostInit` tag can be used during construction for all value types
 - ▶ This tag is **mandatory** when `value_type` is a View
 - ▶ Requires that UnorderedMap is allocated in host-accessible memory space

Example:

```
using view_type = Kokkos::View<int*, ExecSpace>;
using map_type = Kokkos::UnorderedMap<int, view_type, HostSpace>;
map_type umap(Kokkos::view_alloc(Kokkos::SequentialHostInit, ""), N);
```

- ▶ Add support for type conversions between `simd` types
- ▶ Make reduction identity of half-precision types return the correct type
 - ▶ Return `[b]half_t` instead of `float` for `reduction_identity<...>` member functions
 - ▶ Concerns the functions: `sum()`, `prod()`, `min()`, and `max()`
- ▶ Improve `Kokkos::Array` consistency and compatibility with STL array
 - ▶ Add `[c]begin()` and `[c]end()` methods
 - ▶ Add missing `noexcept` specifier to some methods

- ▶ Print the commit hash of the embedded dependencies at configure time and when calling `print_configuration`
- ▶ Desul atomics: Use Clang atomic min/max GCC-style builtins
- ▶ Enable running tests on systems with only 2GB of device memory
- ▶ Updated ErrorReporter: adhere to Kokkos naming conventions and change `get_reports` to return values

New execution policy trait: Kokkos::Experimental::StaticBatchSize<N>

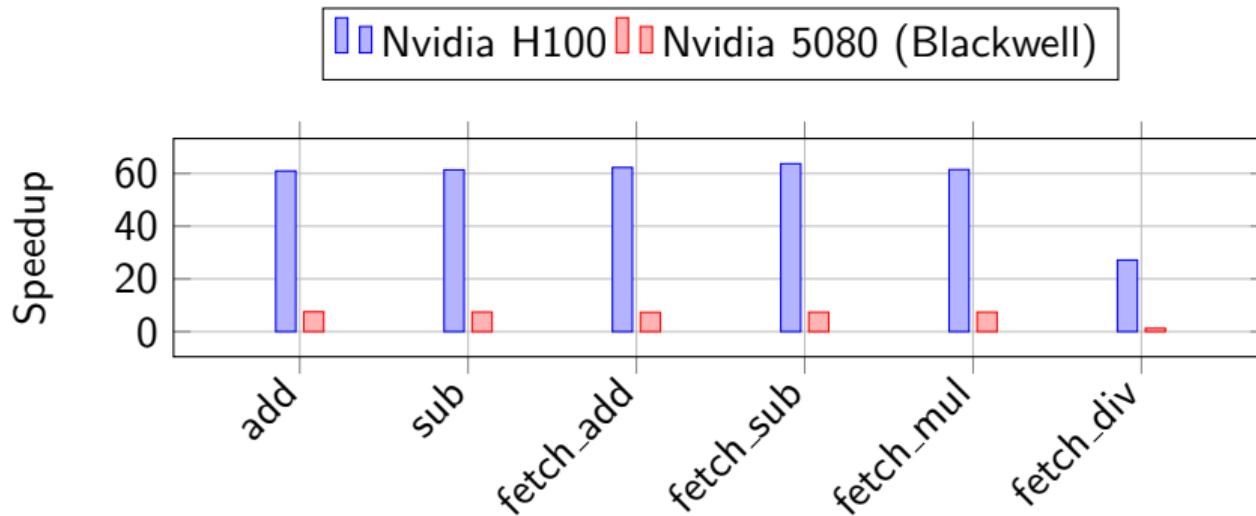
- ▶ Allows users to specify a loop unrolling factor at compile time
- ▶ Only effective for RangePolicy in the CUDA backend for now

Example usage:

```
// Define a batch size of 4
using unroll_4 = Kokkos::Experimental::StaticBatchSize<4>;
// Use in RangePolicy
Kokkos::RangePolicy<ExecSpace, unroll_4> policy(0, N);
// Execute parallel_for with static batch size
Kokkos::parallel_for(policy, KOKKOS_LAMBDA(const int i) {
    ...
});
```

Backend Updates

CUDA



- ▶ Atomics on 10^8 Kokkos::complex<double> **without** contention
 - ▶ Speedup $\approx 60x$ on H100 and $\approx 7x$ on RTX5080.
 - ▶ Same performance for int128 and Kokkos::complex<double>.
 - ▶ Division more costly, thus less effect of atomic CAS.

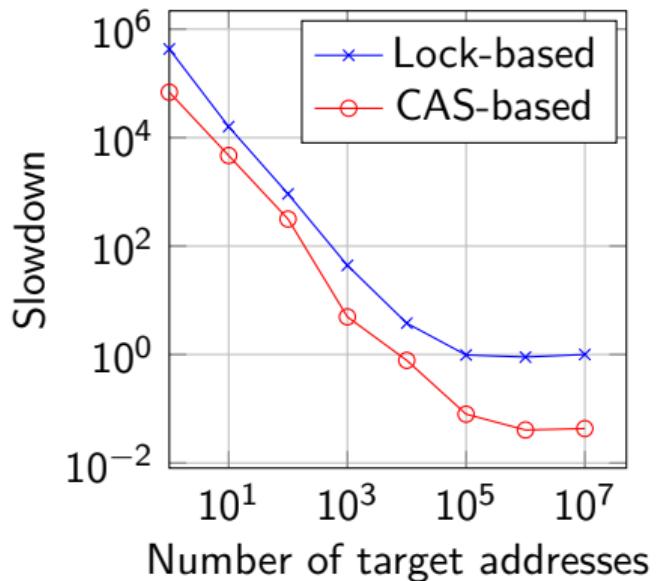


Figure: atomic_add of Kokkos::complex<double> with 10^8 workers

- ▶ Effectiveness of CAS-based atomics reduces similar to Lock-based atomics at high contention.

- ▶ Allows to launch kernels with up to 32kB of arguments for kernels. Previously it was 4kB.
- ▶ Enables us to side-step the "Constant Cache" launch mechanism in Kokkos.
- ▶ Effects functors in the 4kB to 32kB size range. No effect on smaller or larger functors.
- ▶ This changes the synchronization behavior for functors in this range, due to elimination of an implicit necessary synchronization on constant cache buffer use.
- ▶ Does not apply to using Clang as CUDA compiler, nor for GPUs older than Volta (i.e. Compute Capabilities lower than 7).

SYCL

- ▶ SYCL now uses an unsigned integer type as `size_type`.
- ▶ Now unsigned integer type across all backends.

OpenMPTarget

- ▶ Previously parallel_scan with a RangePolicy needed to start at index 0
- ▶ Now any starting index smaller than the end index is supported.

We decided to remove OpenMPTarget in an upcoming release!

- ▶ Never reached feature parity.
- ▶ Lower performance than native backends (CUDA, HIP, SYCL)
- ▶ Practically no users.
- ▶ Little interest in support by any institution.

OpenACC

- ▶ Previously parallel_scan with a RangePolicy needed to start at index 0
- ▶ Now any starting index smaller than the end index is supported.

- ▶ OpenACC now supports the Kokkos_Random algorithms API.
- ▶ Can be inefficient if the actual team size is different from the default team size.

- ▶ OpenACC now supports the `partition_space` API.
- ▶ Execution space instances created by `partition_space` will use OpenACC async IDs in a reserved range (from 64 to 191), which are assigned in a round-robin manner.

- ▶ OpenACC now supports custom scalar reduction with `parallel_reduce` and `RangePolicy`.
- ▶ Supports both built-in reducers with custom scalar types, and custom reducers with custom scalar types.

HIP

- ▶ Fix a performance regression introduced in 4.6 when using lightweight kernel (`Experimental::WorkItemProperty::HintLightWeight`) in `parallel_reduce`
- ▶ Prefer smaller block sizes for `parallel_for` when the requested parallelism is less than the available concurrency
- ▶ Use atomic builtins for `atomic_fetch_min/max` with floating point types instead of our own implementation
- ▶ Add support for Navi4 architecture (Radeon AI PRO R9700, Radeon RX 9070 XT)

- ▶ Avoid using ROCm 7.1 if possible: **you may get incorrect results**
- ▶ On MI100 and MI200 series, use
 - DKokkos_ENABLE_IMPL_HIP_MALLOC_ASYNC=OFF
- ▶ On MI300 series, we cannot compile the testsuite yet. Very likely that you will also need to use -DKokkos_ENABLE_IMPL_HIP_MALLOC_ASYNC=OFF

Build System Updates

- ▶ New (unsurprising) defaults
 - ▶ Kokkos_ENABLE_DEPRECATED_CODE_4=OFF
 - ▶ Kokkos_ENABLE_DEPRECATED_CODE_5=ON
- ▶ CMake Requirements
 - ▶ CMake 3.22 is the new minimum
 - ▶ CMake 3.25.2 is the minimum for CUDA CMake language support

- ▶ Kokkos now offers support for multiple CMake languages in one build
 - ▶ Allow a single Kokkos installation to be used with the CMake languages CXX and CUDA/HIP.
 - ▶ Kokkos sets the flags for all CMake languages that work with the selected backend.
 - ▶ Rules for `kokkos_launch_compiler` stay unchanged: Used for CUDA backend **without** `nvcc_wrapper` as `CXX_COMPILER`
 - ▶ Enable via `Kokkos_ENABLE_MULTIPLE_CMAKE_LANGUAGES=ON`

CMake

```
find_package(Kokkos)

# use Kokkos in a CUDA language project
add_library(language_lib library.cu)
target_link_libraries(language_lib PUBLIC Kokkos::kokkos)

#use Kokkos in a CXX project
add_library(cxx_lib library.cpp)
target_link_libraries(cxx_lib PUBLIC Kokkos::kokkos)

#use both in an executable
add_executable(example cmake_example.cpp)
target_link_libraries(example PUBLIC language_lib cxx_lib)
```

- ▶ Look into examples/build_cmake_installed_multilanguage

- ▶ Experimental support for shared library builds in Windows
 - ▶ Kokkos now supports `BUILD_SHARED_LIBS=ON` in Windows builds
 - ▶ Includes `WINDOWS_EXPORT_ALL_SYMBOLS=ON` for all Kokkos-core libraries

Deprecations

New Kokkos_ENABLE_DEPRECATED_CODE_5 (= ON by default)

- ▶ Deprecate KOKKOS_ATTRIBUTE_NODISCARD macro
- ▶ Deprecate Owning, ObservingRawPtr aliases
- ▶ Deprecate Random_XorShift{64,1024}_Pool::init
- ▶ Deprecate basic_simd::{const_}where_expression
- ▶ Deprecate support for using nested OpenMP parallel regions without nested OpenMP enabled
- ▶ Deprecate creating Kokkos::OpenMP instances inside OpenMP parallel regions

Various type alias deprecations and replacements in View

- ▶ HostMirror → host_mirror_type
- ▶ scalar_array_type → data_type
- ▶ const_scalar_array_type → const_data_type
- ▶ non_const_scalar_array_type → non_const_data_type
- ▶ array_type → type
- ▶ DynamicView::array_type → DynamicView::uniform_type

Various deprecations and replacements in ErrorReporter

- ▶ `getCapacity()` → `capacity()`
- ▶ `getNumReports()` → `num_reports()`
- ▶ `getNumReportAttempts()` → `num_report_attempts()`
- ▶ `getReports()` → `get_reports()`

```
// before
std::vector<int> reporters;
std::vector<report_type> reports;
error_reporter.getReports(reporters, reports);

// now uses structured bindings
auto [reporters, reports] = error_reporter.get_reports();
```

Note

- ▶ `Kokkos_ENABLE_DEPRECATED_CODE_4 = ON` → `OFF`

Breaking Changes

Makefiles

- ▶ Deprecated in 4.6, dropped in 5.0.
- ▶ Removes the burden of maintaining two build systems

Kepler architecture (Nvidia GPU)

- ▶ sm_3x not supported since CUDA 12.
- ▶ Dropping support for Kokkos_ARCH_KEPLER30/32/35/37.

Google Benchmark

- ▶ v1.8.3 is the new minimum.

Task-based parallelism

- ▶ Rarely used. Deprecated in 4.5. Now, dropped.
- ▶ No more: TaskSingle, task_spawn, host_spawn, BasicFuture, etc.

Parallel dispatch: parallel_for/scan/reduce

- ▶ Dispatch only within the Kokkos execution environment.
- ▶ Abort if used before initialize() or after finalize().

Execution space chosen by default in Kokkos::Graph::submit()

- ▶ Used to be the execution space specified during Graph construction.
- ▶ Now, defaults to Kokkos::DefaultExecutionSpace.
- ▶ Matches semantics of other parallel constructs in Kokkos.

Default SIMD vector width

- ▶ From now, choose the largest SIMD vector width available for type T.
- ▶ Example (with AVX512): Kokkos::Experimental::simd<float> maps to:

```
before:  
    Kokkos::Experimental::basic_simd<float, avx512_fixed_size<8>>  
after:  
    Kokkos::Experimental::basic_simd<float, avx512_fixed_size<16>>
```

- ▶ Specifying desired vector width still available:

```
Kokkos::Experimental::simd<float, 8>
```

Memory trait alias Kokkos::MemoryRandomAccess

- ▶ Also meant Kokkos::Unmanaged memory trait. Now, fixed.
- ▶ Older meaning supported via Kokkos_ENABLE_DEPRECATED_CODE_4.

Bug Fixes

- ▶ Track modification for resize only if DualView is not using a single device
- ▶ Fix MSVC floating-point value does not fit in required floating-point type warning from `reduction_identity`
- ▶ Properly delete Timer copy constructor and copy assignment operators
- ▶ Fix RISC-V support (compiler check at configuration time and missing semicolons at compile time)
- ▶ Corrected `bit_width` return type to be `int` instead of `T` to align with the standard library
- ▶ Work around a performance regression related to index computation in the `mdspan-based View`

- ▶ Launch work graph on the instance in the execution policy
- ▶ Fix incorrect results from `parallel_reduce` with `LaunchBounds` values smaller than 32

- ▶ Fix partition_space abort when requesting more partitions than available threads

- ▶ Fix a configure-time failure when SVE is enabled and unit tests are disabled

How to Get Your Fixes and Features into Kokkos

- ▶ Fork the Kokkos repo (<https://github.com/kokkos/kokkos>)
- ▶ Make topic branch from *develop* for your code
- ▶ Add tests for your code
- ▶ Create a pull request (PR) on the main project *develop*
- ▶ Update the documentation (<https://github.com/kokkos/kokkos-core-wiki>) if your code changes the API
- ▶ Get in touch if you have any question (<https://kokkosteam.slack.com>)